

# Statistical Language Models

Week 7.2

Textmining in R, chapter 8

# Perplexity

- Training dataset used to build LDA model
- Validation dataset used to assess the (model selection) number of clusters topics

# From hold out data

- Great resource: Section 2 of
- vignette("topicmodels")
- Geometric mean per word likelihood:

- $$Perplexity(Doc_{test}) = \exp \left\{ - \frac{\sum_{doc} \log(p(w \mid \theta, \phi))}{\sum_{doc} \{\#words \text{ is in doc} \}} \right\}$$

- lower perplexity score indicates better generalization performance

# Perplexity

- Measure is how likely new data is given the LDA model that was learned.
- Assumptions: standard to training /testing
- Perplexity is not well correlated with human perception of quality by interpretability of topics

# JavaScript Object Notation (JSON) formats

- JSON data comes in key/value pairs
- Think of it as named data frame columns or named vectors.
- Our course schedule might be have several key/value pairs:

"startTime":"13:05", "days":"Tu/Th",

"endTime":"14:25",

"roomNumber":"3224","building":"Richcraft Hall"

# JavaScript Object Notation (JSON) formats

- Data can be
  - Numbers: 1, 2, 3.1415
  - Strings: "Text in double quotes"
  - Boolean: TRUE
  - Array ["ordered", "comma separated", "enclosed in square brackets", "any data type inside"]
  - Object {unordered, comma separated, collection of key: value pairs in curly brackets, any data types}

# JavaScript Object Notation (JSON) formats

- Our course schedule might be have several  
“key”:value pairs:

"startTime":"13:05", "days":"Tue/Thu",

"endTime":"2:25",

"roomNumber":"3224","building":"Richcraft Hall"

# JavaScript Object Notation (JSON) formats

- Our course info could be split into a hierarchical data structure with the top levels:
- CourseSchedule
- Course Instructor
- Grading
- Course Info
- Text



# JavaScript Object Notation (JSON) formats

- Our course info could be split into a hierarchical data structure with the top levels:
- CourseSchedule [start time, end time, day, room, ...]
- Course Instructor [name, office, email, phone,...]
- Grading [assignments, midterms, final,...]
- Course Info [pre-req, delivery method, title,description,...]
- Text [required, recommended,...]

# Some course info in JSON:

```
“courseSchedule”:[{"startTime":"13:05",  
"startDate":"Tue Jan 07 2020", "roomNumber":"7618",  
"days":"Mo","endDate":"Thu Apr 07  
2020","endTime":"14:20", "isExam":false,  
“roomNumber":"3224","building":"Richcraft Hall"}],
```

# Big JSON file

- Can we find datasets that are related? Are there clusters of similar datasets? What does NASA research?
- <https://data.nasa.gov/data.json>

# Metadata from NASA datasets from their research

- `library(jsonlite)`
- `metadata = fromJSON("https://data.nasa.gov/data.json") # big file`
- `class(metadata)`
- `names(metadata$dataset)`
-

- `class(metadata$dataset$title)`
- `length(metadata$dataset$title)`
- `class(metadata$dataset$keyword)`
- `class(metadata$dataset$description)`

# Tidying up the title data

- #Set up separate nibbles for title, description, and keyword
- library(dplyr)
- nasa\_title = tibble(id = metadata\$dataset\$`identifier`,
- title = metadata\$dataset\$title)
- nasa\_title
-

# Data set descriptions

- `nasa_desc = tibble(id = metadata$dataset`identifier`,`
- `desc = metadata$dataset$description)`
- `nasa_desc %>%`
- `select(desc) %>%`
- `sample_n(5)`
-

# The data keywords (one row for each keyword)

- `library(tidyr)`
- `nasa_keyword = tibble(id = metadata$dataset$`identifier`,`
- `keyword = metadata$dataset$keyword) %>%`
- `unnest(keyword)`
- `nasa_keyword`



# Tokenize title and description

- library(tidytext)
- nasa\_title = nasa\_title %>%
- unnest\_tokens(word, title) %>%
- anti\_join(stop\_words)
- nasa\_title
- ##### stop\_words????  
#####
- nasa\_desc = nasa\_desc %>%
- unnest\_tokens(word, desc)  
%>%
- anti\_join(stop\_words)
- nasa\_desc

# Basic Explorations

- # Find the most common words in the dataset
- `nasa_title %>%`
- `count(word, sort = TRUE)`
- `nasa_desc %>%`
- `count(word, sort = TRUE)`

# Custom stop words

## #arguably include words like "data"

- `my_stopwords = tibble(word = c(as.character(1:10),`
- `"v1", "v1.0", "v03", "l2", "l3", "l4", "v5.2.0",`
- `"v003", "v004", "v005", "v006", "v7"))`
- `nasa_title = nasa_title %>%`
- `anti_join(my_stopwords)`
- `nasa_desc = nasa_desc %>%`
- `anti_join(my_stopwords)`

# Common keywords

- `nasa_keyword %>%`
- `group_by(keyword) %>%`
- `count(sort = TRUE)`

# Networks of Description and Title Words

- #Count the number of pairwise occurrences of words in a title or description
- library(widyr)
- title\_word\_pairs = nasa\_title %>%
- pairwise\_count(word, id, sort = TRUE, upper = FALSE)
- title\_word\_pairs

# Common co-occurring description words

- `desc_word_pairs = nasa_desc %>%`
- `pairwise_count(word, id, sort = TRUE, upper = FALSE)`
- `desc_word_pairs`

# Networks of co-occurring words in titles

- `library(ggplot2)`
- `library(igraph)`
- `library(ggraph)`
- `title_word_pairs %>%`
- `filter(n >= 250) %>%`
- `graph_from_data_frame() %>%`
- `ggraph(layout = "fr") +`
- `geom_edge_link(aes(edge_alpha = n,`  
`edge_width = n), edge_colour = "cyan4")`  
`+`
- `geom_node_point(size = 5) +`
- `geom_node_text(aes(label = name),`  
`repel = TRUE,`
- `point.padding = unit(0.2,`  
`"lines")) +`
- `theme_void()`

# Networks of Co-occurrences in descriptions

- `desc_word_pairs %>%`
- `filter(n >= 1000) %>%`
- `graph_from_data_frame() %>%`
- `ggraph(layout = "fr") +`
- `geom_edge_link(aes(edge_alpha = n,`  
`edge_width = n), edge_colour =`  
`"darkred") +`
- `geom_node_point(size = 5) +`
- `geom_node_text(aes(label = name),`  
`repel = TRUE,`
- `point.padding = unit(0.2,`  
`"lines")) +`
- `theme_void()`



# Networks of co-occurrences in keywords

- `keyword_pairs = nasa_keyword %>%`
- `pairwise_count(keyword, id, sort = TRUE, upper = FALSE)`
- `keyword_pairs`
- `keyword_pairs %>%`
- `filter(n >= 500) %>%`
- `graph_from_data_frame() %>%`
- `ggraph(layout = "fr") +`
- `geom_edge_link(aes(edge_alpha = n, edge_width = n), edge_colour = "royalblue") +`
- `geom_node_point(size = 5) +`
- `geom_node_text(aes(label = name), repel = TRUE,`
- `point.padding = unit(0.2, "lines")) +`
- `theme_void()`

- #Note the clusters; these seem to define redundancy, but counts won't tell us for sure. Co-occurrence isn't always enough
- keyword\_cors = nasa\_keyword %>%
- group\_by(keyword) %>%
- filter(n() >= 50) %>%
- pairwise\_cor(keyword, id, sort = TRUE, upper = FALSE)
- keyword\_cors

# Correlation options

- Pearson Correlation is classic linear relationship for two continuous variables
- Kendall rank correlation measures the strength of dependence for ordered data

- $$\tau = \frac{N_{concordant} - N_{discordant}}{.5N(N - 1)} = \frac{1}{.5N(N - 1)} \sum_{i < j} \text{sgn}(x_i - x_j) \text{sgn}(y_i - y_j)$$

- Concordant = consistent
- Spearman rank correlation for ordinal or continuous variables

- $$\rho = 1 - \frac{6 \sum d_i^2}{N(N^2 - 1)}, \text{ for difference between ranks } d_i \text{ and } N \text{ observations}$$

# Visualize networks of words via correlations

- `keyword_cors %>%`
- `filter(correlation > .6) %>%`
- `graph_from_data_frame() %>%`
- `ggraph(layout = "fr") +`
- `geom_edge_link(aes(edge_alpha = correlation, edge_width = correlation), edge_colour = "royalblue") +`
- `geom_node_point(size = 5) +`
- `geom_node_text(aes(label = name), repel = TRUE,`
- `point.padding = unit(0.2, "lines")) +`
- `theme_void()`

# tf\_idf from descriptions

- `desc_tf_idf = nasa_desc %>%`
- `count(id, word, sort = TRUE) %>%`
- `ungroup() %>%`
- `bind_tf_idf(word, id, n)`
- `desc_tf_idf = full_join(desc_tf_idf, nasa_keyword, by = "id")`
- `desc_tf_idf %>%`
- `arrange(-tf_idf)`
- #many boring words

# Look for high tf\_idf terms within keywords

- `desc_tf_idf %>%`
- `filter(!near(tf, 1)) %>%`
- `filter(keyword %in% c("solar activity", "clouds", "seismology", "astrophysics", "human health", "budget", "climate" )) %>%`
- `arrange(desc(tf_idf)) %>%`
- `group_by(keyword) %>%`
- `distinct(word, keyword, .keep_all = TRUE) %>%`
- `top_n(15, tf_idf) %>%`
- `ungroup() %>%`
- `mutate(word = factor(word, levels = rev(unique(word)))) %>%`
- `ggplot(aes(word, tf_idf, fill = keyword)) +`
- `geom_col(show.legend = FALSE) +`
- `facet_wrap(~keyword, ncol = 3, scales = "free") +`
- `coord_flip() +`
- `labs(title = "Highest tf-idf words in NASA metadata description fields",`
- `caption = "NASA metadata from https://data.nasa.gov/data.json",`
- `x = NULL, y = "tf-idf")`

# LDA

- `my_stop_words = bind_rows(stop_words, tibble(word = c("nbsp", "amp", "gt", "lt", "timesnewromanpsmt", "font", "td", "li", "br", "tr", "quot", "st", "img", "src", "strong", "http", "file", "files", as.character(1:12)), lexicon = rep("custom", 30)))`
- `word_counts = nasa_desc %>%`
- `anti_join(my_stop_words) %>%`
- `count(id, word, sort = TRUE) %>%`
- `ungroup()`
- `word_counts`

# To DTM

- `desc_dtm = word_counts %>%`
- `cast_dtm(id, word, n)`
- `desc_dtm`
- `### desc_dtm = removeSparseTerms(desc_dtm,.95)`



# LDA

- `library(topicmodels)`
- # be aware that running this model is time intensive
- `desc_lda <- LDA(desc_dtm, k = 24, control = list(seed = 1234))`
- `desc_lda`

- `tidy_lda <- tidy(desc_lda)`

- `tidy_lda`

# Their $\beta$ = probability of a term given a topic

- `top_terms <- tidy_lda %>%`
- `group_by(topic) %>%`
- `top_n(10, beta) %>%`
- `ungroup() %>%`
- `arrange(topic, -beta)`
- `top_terms`

# In plots

- `top_terms %>%`
- `mutate(term = reorder_within(term, beta, topic)) %>%`
- `group_by(topic, term) %>%`
- `arrange(desc(beta)) %>%`
- `ungroup() %>%`
- `ggplot(aes(term, beta, fill = as.factor(topic))) +`
- `geom_col(show.legend = FALSE) +`
- `coord_flip() +`
- `scale_x_reordered() +`
- `labs(title = "Top 10 terms in each NASA topic",`
- `x = NULL, y = expression(beta)) +`
- `facet_wrap(~ topic, ncol = 4, scales = "free")`

# Their $\gamma$ = probability of topic within document

## Connecting topic modeling with keywords

- `lda_gamma = tidy(desc_lda, matrix = "gamma")`
- `lda_gamma`
- `lda_gamma = full_join(lda_gamma, nasa_keyword, by = c("document" = "id"))`
- `lda_gamma`

# In plots

- `top_keywords <- lda_gamma %>%`
- `filter(gamma > 0.9) %>%`
- `count(topic, keyword, sort = TRUE)`
- `top_keywords`
- `top_keywords %>%`
- `group_by(topic) %>%`
- `top_n(5, n) %>%`
- `ungroup %>%`
- `mutate(keyword = reorder_within(keyword, n, topic)) %>%`
- `ggplot(aes(keyword, n, fill = as.factor(topic)))`  
+
- `geom_col(show.legend = FALSE) +`
- `labs(title = "Top keywords for each topic",`  
`x = NULL, y = "Number of documents") +`
- `coord_flip() +`
- `scale_x_reordered() +`
- `facet_wrap(~ topic, ncol = 4, scales = "free")`

# Supervised LDA

- <https://papers.neurips.cc/paper/3328-supervised-topic-models.pdf>

- Fit LDA so as to best estimate  $Y$

