

# Statistical Language Models

Week 8

# NASA

- NASA LDA with 24 topics
- `library(topicmodels)`
- # be aware that running this model is time intensive
- `desc_lda = LDA(desc_dtm, k = 24, control = list(seed = 1234))`
- `desc_lda`

- `tidy_lda <- tidy(desc_lda)`

- `tidy_lda`

# Their $\beta$ = probability of a term given a topic

- `top_terms <- tidy_lda %>%`
- `group_by(topic) %>%`
- `top_n(10, beta) %>%`
- `ungroup() %>%`
- `arrange(topic, -beta)`
- `top_terms`

# In plots

- `top_terms %>%`
- `mutate(term = reorder_within(term, beta, topic)) %>%`
- `group_by(topic, term) %>%`
- `arrange(desc(beta)) %>%`
- `ungroup() %>%`
- 
- 
- `ggplot(aes(term, beta, fill = as.factor(topic))) +`
- `geom_col(show.legend = FALSE) +`
- `coord_flip() +`
- `scale_x_reordered() +`
- `labs(title = "Top 10 terms in each NASA topic",`
- `x = NULL, y = expression(beta)) +`
- `facet_wrap(~ topic, ncol = 4, scales = "free")`

# Their $\gamma$ = probability of topic within document

## Connecting topic modeling with keywords


- `lda_gamma = tidy(desc_lda, matrix = "gamma")`
- `lda_gamma`
- `lda_gamma = full_join(lda_gamma, nasa_keyword, by = c("document" = "id"))`
- `lda_gamma`

# In plots

- `top_keywords <- lda_gamma %>%`
- `filter(gamma > 0.9) %>%`
- `count(topic, keyword, sort = TRUE)`
- `top_keywords`
- `top_keywords %>%`
- `group_by(topic) %>%`
- `top_n(5, n) %>%`
- `ungroup %>%`
- `mutate(keyword = reorder_within(keyword, n, topic)) %>%`
- `ggplot(aes(keyword, n, fill = as.factor(topic)))`  
+
- `geom_col(show.legend = FALSE) +`
- `labs(title = "Top keywords for each topic",`  
`x = NULL, y = "Number of documents") +`
- `coord_flip() +`
- `scale_x_reordered() +`
- `facet_wrap(~ topic, ncol = 4, scales = "free")`

# Observed and Latent

- Observed: Words in Documents
- Latent: topics for words and documents
- Exploration of Latent space gives interesting insights





# Word2Vec

- Neural Network model to predict a word from those around it
- Produces a numeric latent (embedding) vector space
- <https://arxiv.org/pdf/1301.3781.pdf>

## Matrix Version (not often used in practice)

Based on <https://iksinc.online/tag/continuous-bag-of-words-cbow/>

- Training corpus: {"the dog saw a cat",
- "the dog chased the cat",
- "the cat climbed a tree"}
- Size of vocabulary? (This defined input and output dimensions)
- Define a latent (embedding) of dimension 3 for our vector space

# Model fitting

- Initialize input weight matrix and output weight matrix such that for vocabulary dimension  $V$  and hidden dimension  $H$  we have matrices:
- $\text{INPUT}_{\{V \times H\}}$
- $\text{OUTPUT}_{\{H \times V\}}$
- Ordering the words alphabetically, the input vector representing "dog":
- $X = [0, \dots, 0, 1, 0 \dots 0]$

# Model fitting

- Initialize input weight matrix and output weight matrix such that for vocabulary dimension  $V$  and hidden dimension  $H$  we have matrices:
- $\text{INPUT}_{\{V \times H\}}$
- $\text{OUTPUT}_{\{H \times V\}}$
- Ordering the words alphabetically, the input vector representing "dog":
- $X = [0, \dots, 0, 1, 0 \dots 0]$

- `Vocab = c("a", "cat", "chased", "climbed", "dog", "saw", "the", "tree")`
- `length(Vocab)`
- `V = length(Vocab)`
- `H = 3`
- `INPUT = matrix(rnorm(V*H),nrow=V,ncol = H)`
- `OUTPUT = matrix(rnorm(V*H),nrow=H,ncol = V)`
- `#dog:`
- `X = rep(0,V)`
- `X[Vocab=="dog"]=1`

- #Hidden Layer (transposed)
- $tH = t(X) \% * \% INPUT$
- #Output from the model (no weights trained)
- $Out = tH \% * \% OUTPUT$
- `rbind(Vocab, Out)`

- #Convert to probabilities

- $$P(\text{word} = w \mid \text{word context}) = \frac{\exp[\text{activation}(w)]}{\sum_{v \in \text{Vocab}} \exp[\text{activation}(v)]}$$

- Probs = exp(Out)/sum(exp(Out))
- rbind(Vocab, Out, Probs)

# Model could be optimized

- Hidden layer is a lower dimension vector space.




# Neural Networks

- Massively hierarchical non-parametric models thought to work by magic

# **Mixed** National Institute of Standards and Technology database (MNIST) database


- 28 x 28 pixel images = 784 variables
- 10 different digits to be classified
- A **mix** of American Census Bureau employees and American high school students
- Goal is to classify these digits to human accuracy but faster and cheaper than humans can do it.



- Identify the wine producer from different wine traits.
- 12 wine variables:
- 1) Alcohol 2) Malic acid 3) Ash  
4) Alcalinity of ash 5) Magnesium  
6) Total phenols 7) Flavanoids  
8) Nonflavanoid phenols 9) Proanthocyanins  
10) Color intensity 11) Hue 12) OD280/OD315 of  
diluted wines 13) Proline

# Species identification

- Iris data set is the kindergarten of classification systems
- 3 species of Iris flowers from the Gaspé peninsula



- with 2 input covariates ( $\mathbf{x} = \{\mathbf{A}, \mathbf{B}\}$ ) and  $n=3$  neurons the output function  $o(\mathbf{x})$  is:

$$o(\mathbf{x}) = f \left( w_0 + \sum_{i=1}^n w_i x_i \right)$$

$$o(\mathbf{x}) = f \left( w_0 + \sum_{i=1}^n w_i x_i \right)$$


- $f$  is bounded, monotone, and differentiable.

$$f(u) = \frac{1}{1 + e^{-u}}$$

- Often logistic:

$$\tanh(u) = \frac{e^{2u} - 1}{e^{2u} + 1}$$

- or Hyperbolic Tangent:



# Multiple layers

- Often we build a hierarchical model:

$$o(\mathbf{x}) = f \left( w_0 + \sum_{j=1}^J w_j * f \left( w_{0,j} + \sum_{i=1}^n w_{i,j} x_{i,j} \right) \right)$$

- data  $\longrightarrow$  many parallel neurons  $\longrightarrow$  many parallel neurons  $\longrightarrow \dots \longrightarrow$  predictions
- In the simplest case the inner **f** is the single internal layer and the outer **f** is the output node
- library(neuralnet) uses the same **f** everywhere
- Better tools exist if you don't use R




Figure 1: Example of a neural network with two input neurons (A and B), one output neuron (Y) and one hidden layer consisting of three hidden neurons.

# Models

$$E = \frac{1}{2} \sum_{l=1}^L \sum_{h=1}^H (o_{l,h} - y_{l,h})^2$$

$$o(\mathbf{x}) = f \left( w_0 + \sum_{i=1}^n w_i x_i \right)$$

Evaluation function E

$$f(u) = \frac{1}{1 + e^{-u}}$$

Model output o

Activation function f

- How does this relate to GLMs?



# Model Training

- Fitting criteria, is often Squared Error Loss (i.e. Gaussian likelihood)

$$E = \frac{1}{2} \sum_{l=1}^L \sum_{h=1}^H (o_{l,h} - y_{l,h})^2$$

- Or Cross Entropy (i.e. log Binomial likelihood):

$$E = - \sum_{l=1}^L \sum_{h=1}^H [y_{l,h} \log(o_{l,h}) + (1 - y_{l,h}) \log(1 - o_{l,h})]$$

- for observation  $l$  at output node  $h$

# Model Training

- Optimization is usually gradient based.

$$\frac{\partial E}{\partial w} = 0|_{w=\hat{w}}$$

- Both **E** and **f** are differentiable, so gradients are analytic and often auto-differentiated (call this use of the chain rule *back propagation*)

$$\frac{\partial E}{\partial w} = \frac{dE}{do} \frac{do}{df} \frac{df}{dw}$$

- Often use CG variant or (random) subsets of dimensions to optimize at a time

# Model Training

- Usually BFGS, CG, ... define a step based on curvature.

- Back propagation uses learning rate  $\eta$

$$w_k^{(t+1)} = w_k^{(t)} - \eta_k^{(t)} \left( \frac{\partial E^{(t)}}{\partial w_k^{(t)}} \right)$$

- iteration  $t$  and weight  $k$

# Improvements


- Shutting off some weights via thresholding: sparse auto-encoders
- Imposing known structure: Convolution Neural Nets
- (much) better software

# Real ML software

- Tensorflow, Theano, H2O, Caffee..., etc subdivide the neural net pieces to different GPU cores.
- Model Evaluation and optimization occur very quickly using parallel disjoint model segments and gradients.
- Several orders of magnitude speedup over R is typical

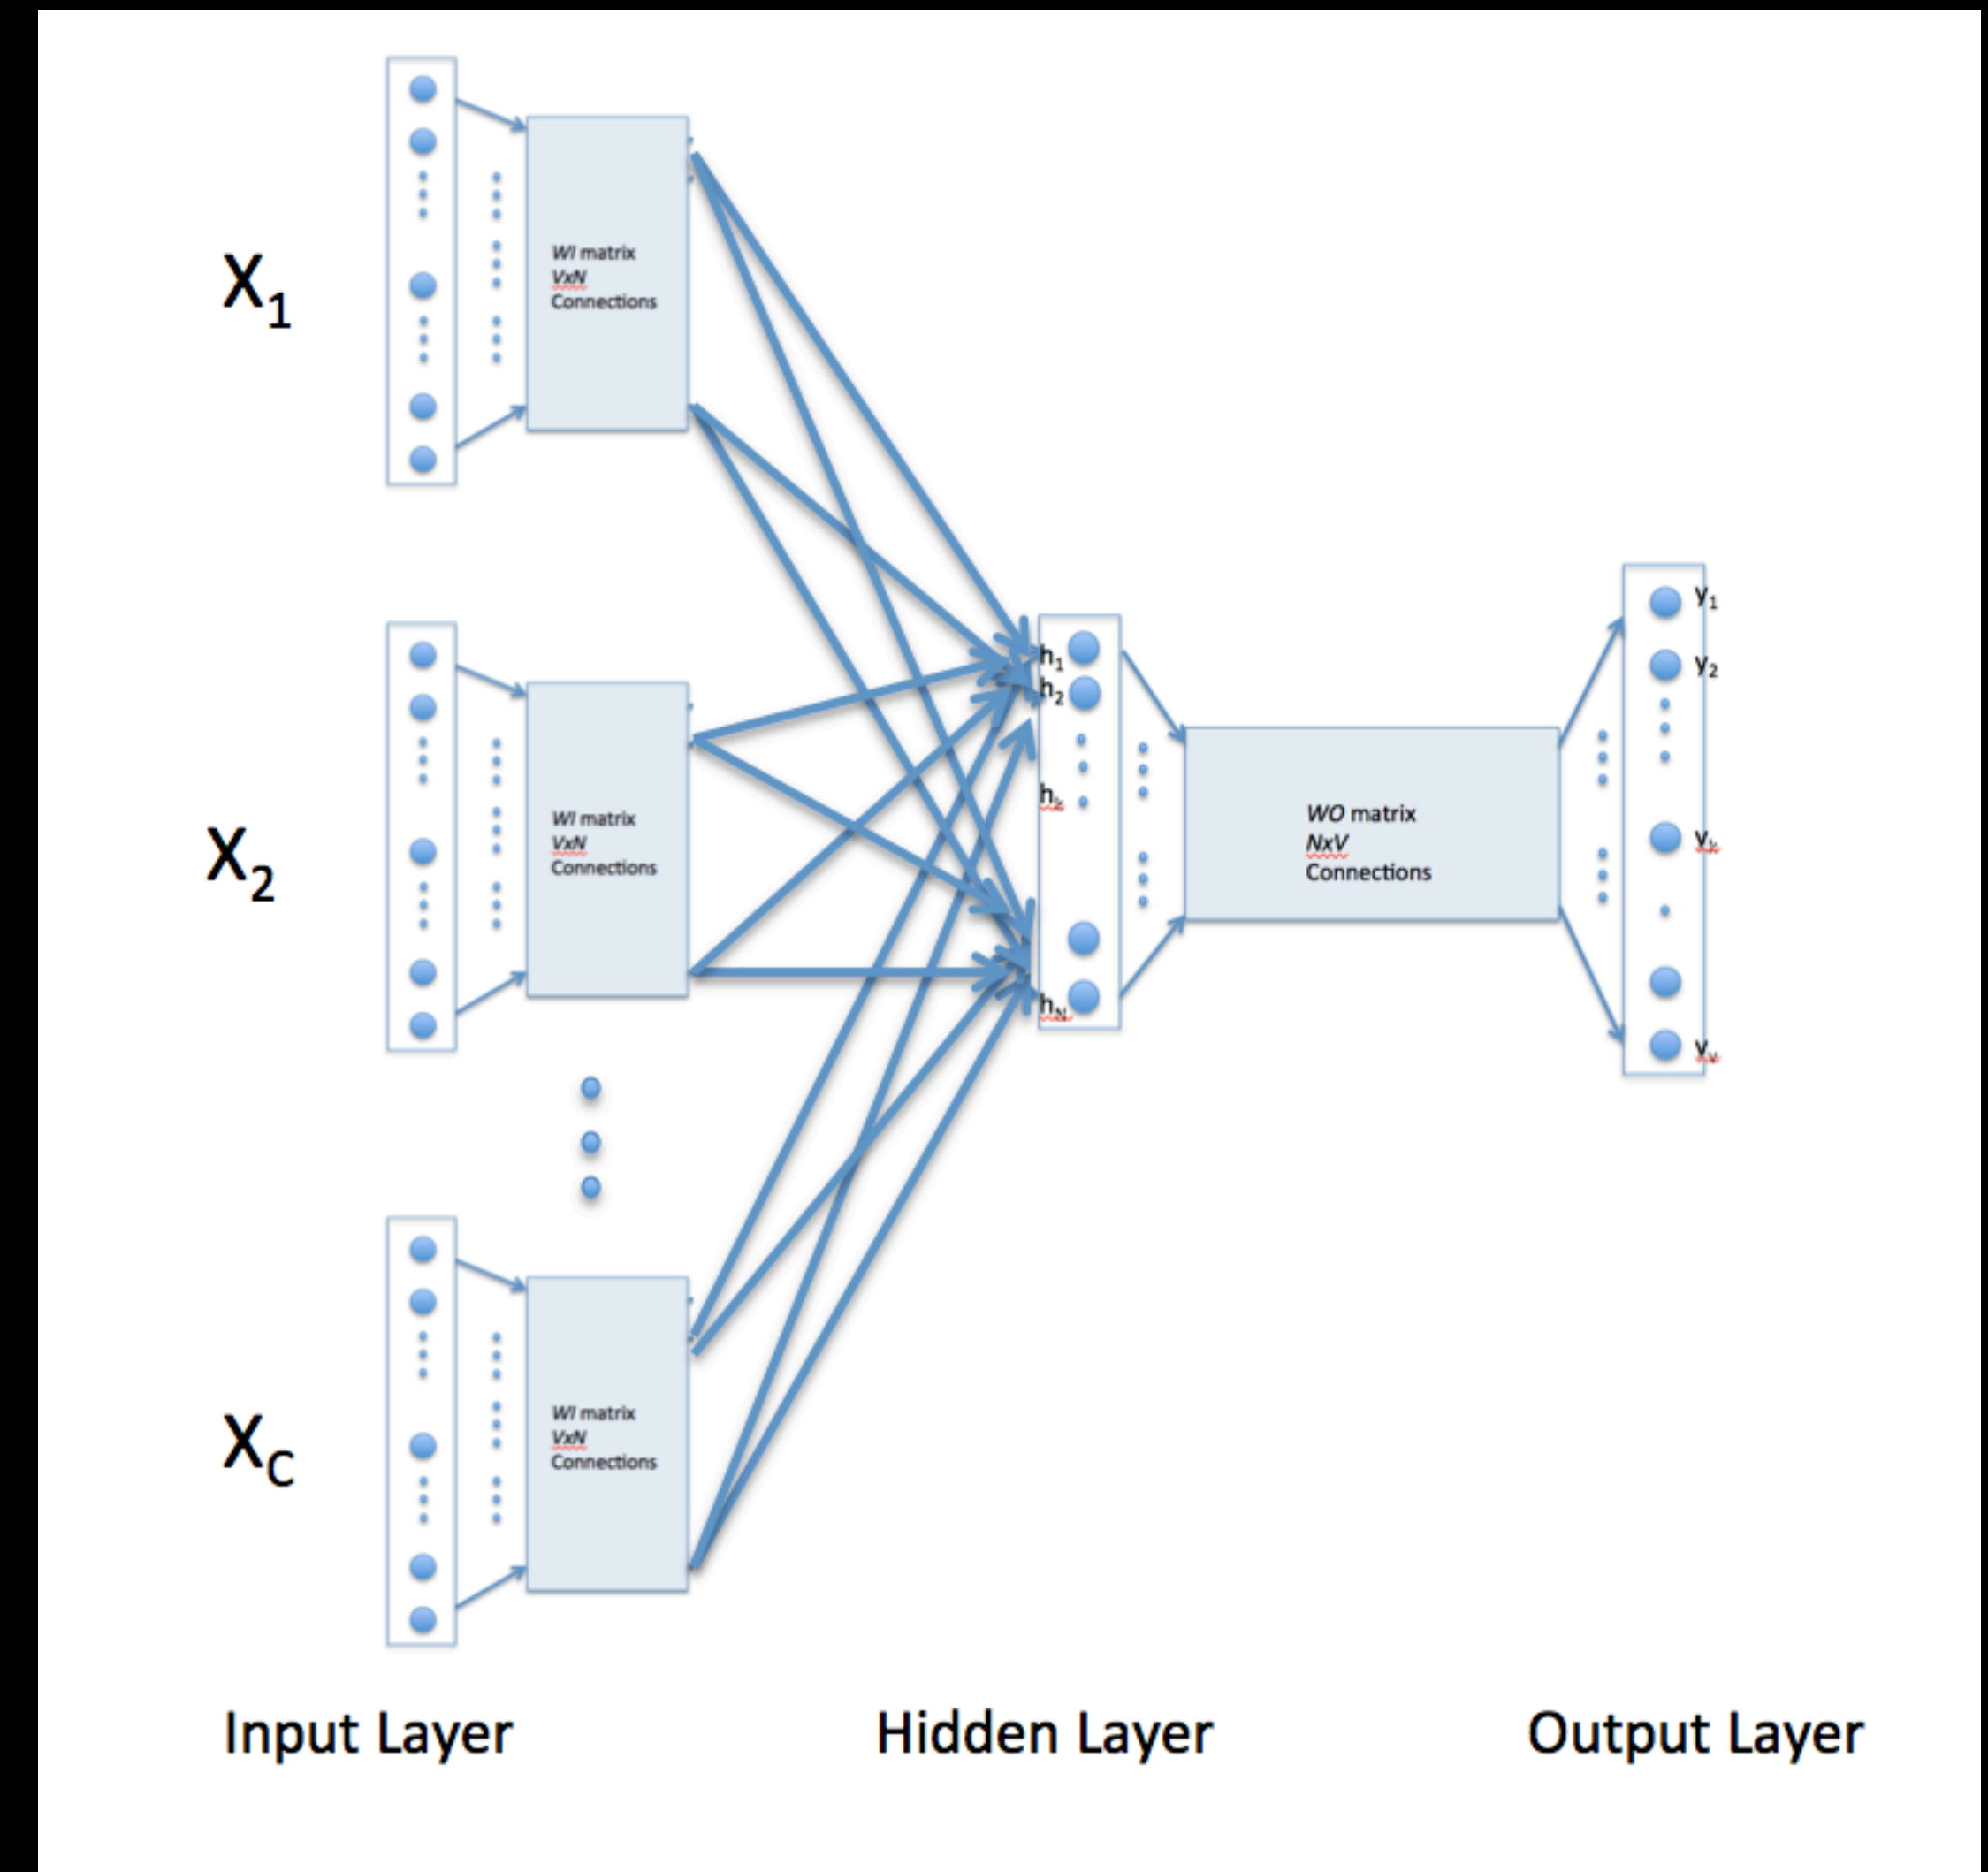
# Word2Vec

- Goal: predict missing word(t)
- Input: {word(t-1), word(t-2), word(t+1), word(t+2)}
- aka: "Continuous Bag of Words" model for predicting centre word from context; BUT word order does not matter




# Word2Vec

- Goal: predict missing word( $t$ )
- Input: {word( $t-1$ ), word( $t-2$ ), word( $t+1$ ), word( $t+2$ )}
- Hidden layer: numerically combines weights and one-hot-encoded covariates



# Word2Vec

- Goal: predict missing word(t)
- Input: {word(t-1), word(t-2), word(t+1), word(t+2)}
- Hidden layer: numerically combines weights and one-hot-encoded covariates
- Output: best word(t)
- Note input dimension and output dimension are the same. Internal dimension (hidden) determines the dimension of the embedding space





# Word2Vec

- `library(devtools)`
- `library(httr)`
- `library(tm)`
- `install_github("bmschmidt/wordVectors")` # yup install from GitHub
- `library(wordVectors)`
- `vignette("introduction", "wordVectors")`

- Main tool is:
- `train_word2vec(InputFileName,OutputFileName,vectors=LatentDimension, threads=CPUCores>window=ContextWindow,iter=ObviouslyMoreIsBetter ButSlower)`
- `train_word2vec("cookbooks.txt","cookbook_vectors.bin",vectors=200,threads=4>window=12,iter=5)`